

A Robust Approach to Automatic Groove Identification in 3D Bullet Land Scans

Kiegan Rice, M.S., Heike Hofmann, Ph.D., Ulrike Genschel, Ph.D.

1 Background

Historically, forensic firearms examiners have analyzed bullet striations through a process of visual feature comparison. Examiners compare striation marks impressed on the lands of a bullet between a bullet fired from a known barrel to a questioned bullet to investigate whether the two bullets were propelled through the same gun barrel.

These visual analyses, based on the concept of Quantitative Consecutively Matching Striae (QCMS), are one of several feature comparison methods whose scientific foundations were questioned by a 2009 report by the National Research Council on Identifying the Needs of the Forensic Sciences Community (National Research Council 2009).

Following that 2009 report, researchers began more intensely studying the validity of feature comparison methods as well as investigating the feasibility of developing image-analysis algorithms to complete automated, quantitative analyses. The main technological development that has created a pathway for image-analysis techniques is the introduction of high resolution 3D scanning technology to the field of forensic science.

3D scanning technology not only allows for preservation of current and historical evidence in digital format, it also provides extremely detailed representations of forensically relevant portions of fired bullets. In recent years, this technology has been applied to the collection of topological images of both bullet lands and breech faces (e.g. De Kinder et al. 1998; De Kinder and Bonifanti 1999; Bachrach 2002). These 3D data have since been used in the development of several methods of varying complexity for automated comparison of land engraved areas (e.g. Ma et al. 2004; Chu et al. 2010; Chu et al. 2013; Hare, Hofmann, and Carriquiry 2017).

Criticisms of firearms examination in recent years have focused on foundational validity and reliability. This doubly enforces the need for automated algorithms to undergo careful study and validation before they can be reasonably implemented to assist forensic firearms examiners. This development process includes data pre-processing methods that ensure the correct data are being used in automated methods.

Data pre-processing is not usually considered a significant barrier for most research endeavors. However, the nature of the 3D scanning process for land engraved areas (LEAs) introduces a challenging data pre-processing problem. To guarantee capture of an entire land engraved area, scanning across the object must begin and end in the neighboring groove engraved areas (GEAs). This ensures the maximal amount of land surface area can be utilized in image-analysis methods, which will provide the most reliable feature generation and more robust results. This extraneous data collection, while necessary, dictates the most significant step in data pre-processing: correctly identifying between data from LEAs and GEAs.

Dealing with these two areas separately is crucial to ensure accuracy and precision in subsequent processing steps. Removal of data from groove engraved areas significantly reduces the possibility of misidentification of the characteristics used in automated comparisons.

Distinguishing between land and groove engraved areas is a problem at which human vision excels, but it is quite challenging for automatic procedures due to the nature of the data collected: the bullet curvature presents the main structure in the data, but the abrupt change between land and groove engraved areas introduces a competing structure. This overwhelms standard statistical modeling techniques. An early solution based on data smoothing falls prey to misidentification of deep striae as shoulder locations. The following work describes a better solution to this pre-processing problem based on robust statistical methods. In order to distinguish between these areas, we aim to identify “shoulder locations”, the locations at which the LEA ends and the GEAs begin.

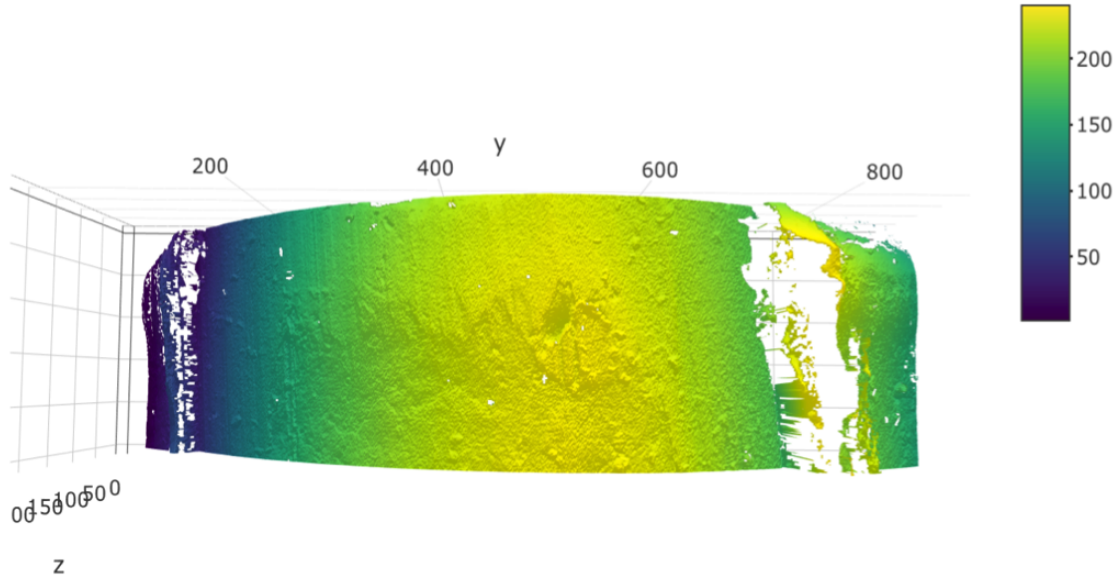


Figure 1: Visualization of 3D data collected through high resolution scanning of a land engraved area. Striations on the surface of the object can be seen by viewing this data from “above”, as presented here.

2 Data Source

The data used in this paper are high resolution 3D scans of bullet land engraved areas. The scanned bullets come from Hamby Set 44 (Hamby, Brundage, and Thorpe 2009). They consist of 35 total bullets from a set of 10 consecutively rifled Ruger P85 barrels. These LEAs were scanned at Iowa State University’s High Resolution Microscopy Facility, and are stored in 3D format as x3p files. **(Add more info about x3p files here??)**. The gathered data are at a resolution of .645 microns per pixel . Physically, each land is approximately 2 millimeters in width, resulting in data structures that can contain up to **this many** individual data points. The 35 total bullets with 6 lands per bullet result in 210 individual lands.

Image-analysis algorithms, while flexible enough to focus on a variety of patterns in the data, should mainly focus on comparison of striation marks and related characteristics that can be calculated. This focus addresses two concerns associated with introducing an automated approach. First, it ensures physical interpretability of characteristics that are calculated from the gathered data. Further, researchers are able to directly compare the visual process examiners use to an automated method which is rooted in the same principles; for example, when a data-based QCMS measure is calculated as part of the algorithm. This striation-focused approach suggests mainly utilizing horizontal slides of the 3D scan, called crosscuts, that capture the striation pattern horizontally across the surface, as seen in Figure 3.

The final data consists of 2D crosscuts gathered from 3D imaging. The height values in the crosscuts were averaged over 10 separate crosscuts spaced out along the 3D image. This ensures predicted locations will be relatively applicable across the depth of the bullet. Two scans were removed from consideration due to data quality concerns, leaving 208 individual crosscuts which function as the dataset.

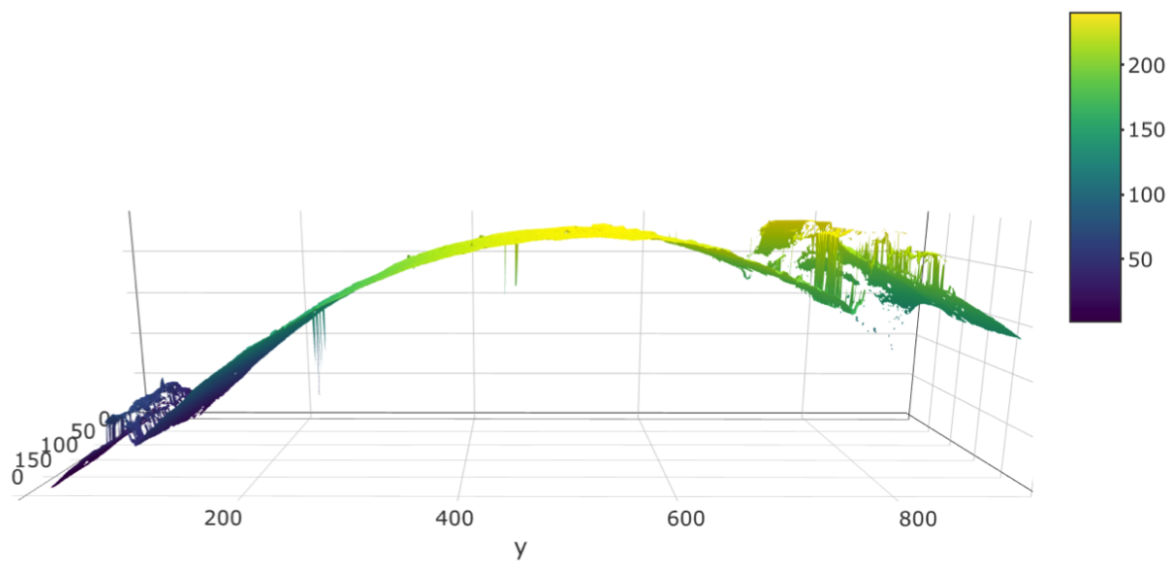


Figure 2: Alternate view of 3D data collected through high resolution scanning of a land engraved area. The globally curved structure of the bullet, as well as portions of the groove engraved area, can be seen clearly when viewing this data from the “side”, as presented here.

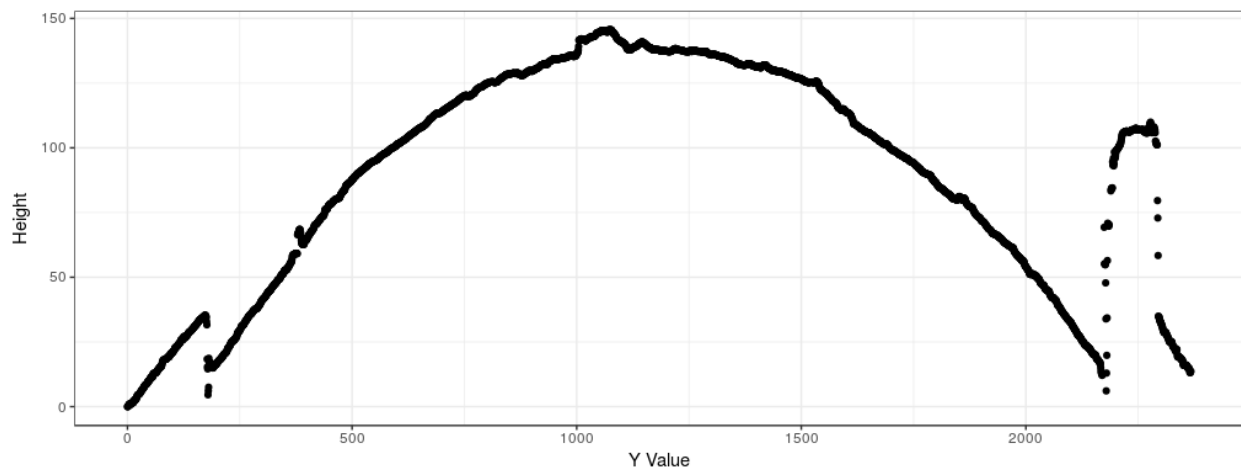


Figure 3: Single crosscut of 3D bullet land data. The main data structure, located in the center, is comprised of the land engraved area. The groove engraved areas occur to the left and right sides of the crosscut.

3 Methodology

The nature of the data structure is such that much of the variability found within the data is due to the global structure of the physical object; that is, the curve of a bullet. Since the ultimate goal is to identify where the global data structure changes (where the shoulder location is), methods need to be able to separate out the LEA structure from the GEA structure. This can most effectively be accomplished by fitting a line to the curve of the bullet and analyzing the pattern of deviations from that curve. That is, fitting a statistical model to the curve of the bullet and examining the residual values.

Due to two competing structures in the data, the ideal statistical model is one which treats the secondary structure of the GEA as outlying data and fits the curve of the LEA alone. We examine two methods for fitting the LEA structure that are based on robust statistical methods. While the two approaches differ in methodology, they are both rooted in the ability to mitigate undue influence caused by outlying data.

3.1 Robust Linear Models

Simple linear regression models assume a one-to-one relationship between the location of a data point on the X axis and the height value on the Y axis. Due to the curved nature of the bullet, it is natural to apply a quadratic linear model to get a curve rather than a straight line fitted to the data. Linear models are based on minimizing the squared distance between each data point and a fitted line. This means that if there are data points in unusual places, a linear model will fit a line that is pulled towards outlying data in order to minimize the overall sum of squared distances. In this particular data environment, the fitted curve can be easily influenced by GEA data, as is seen in Figure 4.

The robust approach under the linear model framework focuses on minimizing the least absolute deviations. This method of minimization is less influenced by possibly large outlying values present in the GEA data. Due to the overwhelming majority of data being from the LEA structure, minimizing the least absolute deviations will favor fitting the LEA structure and allowing GEA data to have large residual values. This is preferable to the traditional linear model, which corrects for the presence of GEA data by compromising between the two competing structures, and does not fit either structure accurately. A striking example of the difference in results from these two model frameworks is seen in Figure 4.

Once a model has been fit for each crosscut, residual values are calculated. The model, having fit the structure of the LEA, results in small residuals scattered around zero in the LEA zone, and larger, mostly positive residuals in the GEA zones. Thus, the magnitudes of the residuals themselves can serve as one indicator of whether a data point is part of the LEA structure or the GEA structure. A cutoff value for the magnitude to separate the residuals from the two zones can be employed. Investigating the median absolute residual (MAR) value, a cutoff that works well to distinguish between GEA and LEA data residuals on the Hamby set 44 is $4 \times \text{MAR}$. Any residual value larger than 4 times the median absolute residual value can be seen as a “large” residual.

Shoulder location predictions are calculated for each crosscut in the following manner:

1. Fit a robust linear model of order 2 (i.e., quadratic) to the averaged crosscut. This is fit using the default methods of the ‘rlm’ function in the ‘MASS’ package in R.
2. Calculate a residual value for each data point on the crosscut.
3. Calculate the median absolute residual (MAR) for the crosscut.
4. Remove all data points on the crosscut whose absolute residual value is greater than $4 \times \text{MAR}$.
5. Find the range of the remaining Y values - these are the predicted shoulder locations for that crosscut.




Figure 4: Example of a quadratic linear model fit (orange) compared to a robust quadratic linear model fit (blue) to a single crosscut. The robust model is able to more effectively capture the curved structure of the LEA without being influenced by the GEA.




Figure 5: Residual values resulting from a robust quadratic linear model. Data points that exist in the GEA structure are clearly separated from the LEA structure.

3.2 Robust LOESS

Locally weighted regression, known as LOESS, is an approach that is not restricted by the need for perfect quadratic curvature. This is advantageous when working with bullets, as it is unrealistic to expect a flawless circular shape to remain after the bullet has been subjected to the forces of a gun barrel and striae have been impressed upon it.

LOESS fits many models to small subsets of the data and combines them into one non-parametric fit of the data, rather than focusing on the overall structure of the data. This allows for greater flexibility. However, it also means that traditional LOESS models are affected by GEA structures in a much more unpredictable manner. A model fitted on a subset of data that mainly falls in the GEA structure will look vastly different than another model fit with data from the LEA. This results in a combined prediction that misrepresents much of the data near one or both shoulder locations (see Figure 6).

Just as the nature of LOESS models differs from traditional linear models, the robust approach must differ too. Robust LOESS utilizes an iterative process focused on re-weighting (see Cleveland 1979). First, an initial LOESS fit is created. This is followed by a step which gives smaller weights to data points with high residual values, and a subsequent LOESS fit with new weights applied. The down-weighting of values with high residual values slowly reduces the influence of a secondary structure within the data; here, the GEA data. This iterative process results in a non-parametric fit to the LEA structure that treats GEA data as less important, which is desirable in this context.

While robust LOESS methods are more flexible than robust linear models, a model that is accurately fit to the LEA structure will result in the same expected residual structure as with robust linear models: positive and negative residuals scattered around zero in the LEA zone, and positive, possibly large residuals in the GEA zones. A similar approach as with the robust linear model cutoff value to distinguish between “large” residual values and reasonable ones; however, because this model is more flexible and fits more closely to the specific data at hand, our cutoff will be lower. A cutoff that performs well on the Hamby set 44 is twice the median absolute residual ($2 \times \text{MAR}$).

Shoulder location predictions are calculated for each crosscut in the following manner:

1. Fit a robust LOESS model with a span of 1 to the averaged crosscut. This is fit using the ‘`locfit.robust`’ function in the ‘`locfit`’ package in R.
2. Calculate a residual value for each data point on the crosscut.
3. Calculate the median absolute residual (MAR) for the crosscut.
4. Remove all data points on the crosscut whose absolute residual value is greater than $2 \times \text{MAR}$.
5. Find the range of the remaining Y values - these are the predicted shoulder locations for that crosscut.

4 Results

In order to assess the accuracy of these predictions, we must also take a unique approach. To calculate a quantitative measure for the overall performance of predictions, “ground truth” shoulder locations were manually identified.

Numerical comparison of predicted and manually identified locations presents a troubling issue; raw distance metrics can misrepresent the true character of a prediction’s accuracy. For example, take a predicted shoulder location that falls 10 data points away from the manually identified shoulder location. This 10-point difference could be caused by noise in the data, missing data points, or simply the miniscule scale of the data. After all, a span of 10 data points represents only 6.45 microns in physical space. Alternatively, a distance of 10 points could actually be 10 points that are part of the groove engraved area, and thus are being incorrectly identified and could later cause problems.

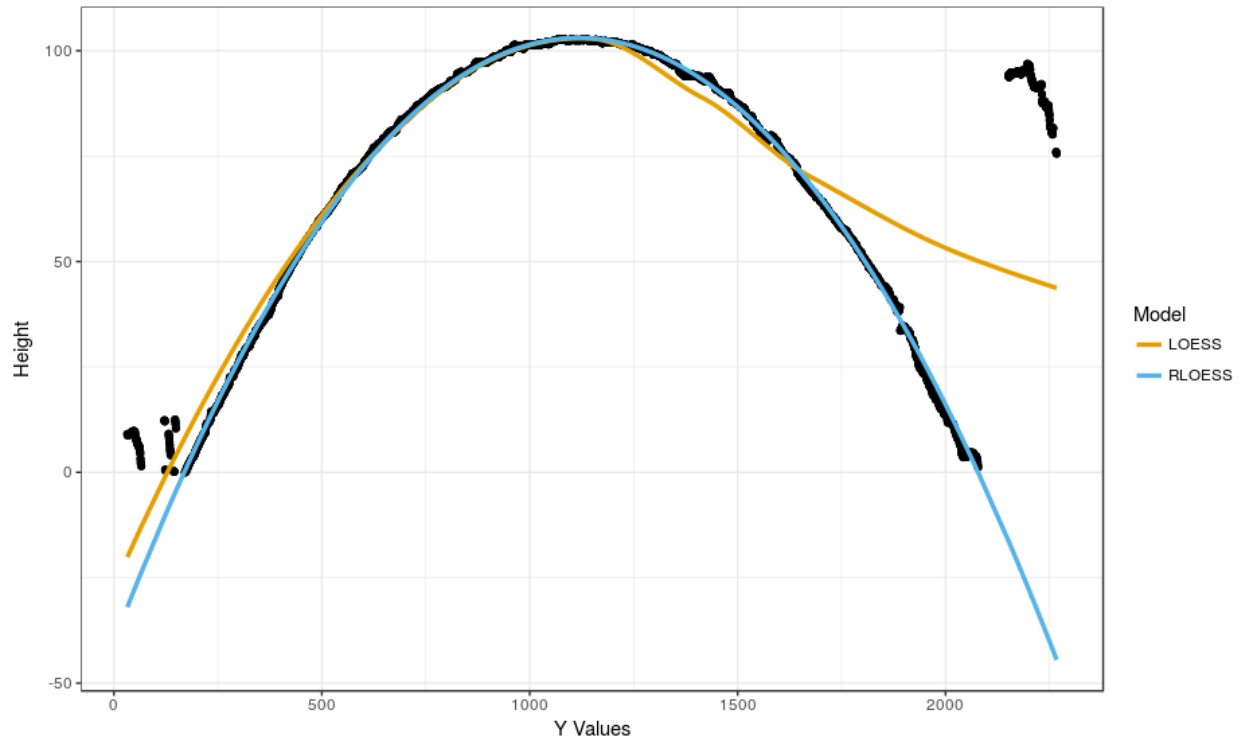


Figure 6: Example of a LOESS model fit (orange) compared to a robust LOESS model fit (blue) to a single crosscut. The robust model is again able to more effectively capture the curved structure of the LEA without being influenced by the GEA.




Figure 7: Residual values resulting from a robust LOESS model. Data points that exist in the GEA structure are clearly separated from the LEA structure. These residuals are very similar to those from the robust quadratic linear model for this particular crosscut.

Thus, a more relevant measure is to investigate the residual values that fall between the predicted and manually identified shoulder location. This penalizes shoulder location predictions that are too far out to the side and leave GEA data in the main structure. Because the robust LOESS most reliably approximates the curved shape of the bullet land due to its flexibility, we want to use residual values resulting from that model to assess final performance of all methods. Residual values from the GEA will not necessarily be uniformly large, but are expected to be positive as their structure and the modeling technique dictates that they would fall above the fitted line from robust LOESS.

Given this assumption, even a 10-point difference can quickly add to a large residual sum if we are dealing with all positive values, as opposed to a 10-point difference within the land engraved area that will be balanced out by the presence of both positive and negative residual values and remain closer to zero.

For this reason, gathering the sum of residuals between the predicted location and the manually identified location is appropriate. This residual sum is referred to as an “inaccuracy score” for which higher values indicate a higher level of inaccuracy. An inaccuracy score was calculated separately for the left and right predictions for each crosscut in the data set.

Of interest are the distributions of these inaccuracy scores across all 208 lands used in the study. A distribution that has a smaller spread and is close to zero is ideal; this suggests many of the predicted shoulder locations are very close to the manually identified locations, and predictions are removing many of the outlying GEA points. A distribution with a wider spread or many high, outlying inaccuracy scores suggests a greater degree of uncertainty and inaccuracy for a particular method.

It is important to note that different results are expected for the left and right shoulder locations. Within Hamby set 44, almost all scans have a well-defined left groove. Left here is defined as visually left on the scan; this is the side the scan begins on, so a well-defined distinction between GEA and LEA is expected. Often, a less clear distinction is seen on the right side of the scan, with sometimes no apparent shoulder location visible. For this reason it is preferable to separate the left and right for visual inspection of results; a method could excel on one side but fall short on another.

Once results are re-run on fresh manual identifications, there may be discussion here of differences between left and right seen in our data.

5 Conclusions

Both the robust linear model and robust LOESS approaches outperform currently implemented solutions based on data smoothers. Of the two, the robust LOESS approach clearly outperforms the robust linear model. This hierarchy of performance is well within expectation given the strength of robust approaches in general as well as the flexibility of LOESS applied to this data type. Robust LOESS also readily handles variation introduced in the process of translating the physical bullet into a 3D object. If there is too much variability in how the bullet is placed relative to the plane of reference (??) on the microscope, crosscuts can have tilted shapes relative to the x-axis which a quadratic linear model would fail to address. In these situations, LOESS excels.

While the cutoff values presented work well on Hamby set 44, additional cross-validation will need to be implemented on a variety of bullet types. Depth of striae, physical size of bullet due to caliber, and non-traditional rifling techniques may require some alterations to this cutoff value. In addition, a study of the effect of implementing a robust LOESS data pre-processing strategy on overall automated image-analysis methods will need to be addressed. Due to increased accuracy of predicted shoulder locations, the authors expect an increase in accuracy in bullet matching algorithms. However, this will need to be validated on a variety of data sets prior to implementation without human intervention in the automated process.




Figure 8: Distribution of inaccuracy scores for data smoothing method, robust linear model method, and robust LOESS method, separated by left and right shoulder locations. A tight distribution with few high values indicates good performance across the LEAs in the data set.

6 References

- Bachrach, Benjamin. 2002. "Development of a 3d-Based Automated Firearms Evidence Comparison System." *Journal of Forensic Sciences* 47 (6): 1253–64.
- Chu, Wei, T. Song, J. Vorburger, J. Yen, S. Ballou, and B. Bacharach. 2010. "Pilot Study of Automated Bullet Signature Identification Based on Topography Measurements and Correlations." *Journal of Forensic Sciences* 55 (2): 341–47.
- Chu, Wei, Robert M Thompson, John Song, and Theodore V Vorburger. 2013. "Automatic Identification of Bullet Signatures Based on Consecutive Matching Striae (Cms) Criteria." *Forensic Science International* 231 (1-3): 137–41.
- Cleveland, William S. 1979. "Robust Locally Weighted Regression and Smoothing Scatterplots." *Journal of the American Statistical Association* 74 (368): 829–36.
- De Kinder, J., and M. Bonifanti. 1999. "Automated Comparison of Bullet Striations Based on 3d Topography." *Forensic Science International* 101: 85–93.
- De Kinder, J., P. Prevot, M. Pirlot, and B. Nys. 1998. "Surface Topology of Bullet Striations: An Innovating Technique." *AFTE Journal* 30 (2): 294–99.
- Hamby, James E., David J. Brundage, and James W. Thorpe. 2009. "The Identification of Bullets Fired from 10 Consecutively Rifled 9mm Ruger Pistol Barrels: A Research Project Involving 507 Participants from 20 Countries." *AFTE Journal* 41 (2): 99–110.
- Hare, Eric, Heike Hofmann, and Alicia Carriquiry. 2017. "Automatic Matching of Bullet Land Impressions." *The Annals of Applied Statistics* 11 (December): 2332–56.
- Loader, Catherine. 2013. "Locfit: Local Regression, Likelihood and Density Estimation." Computer Program. <https://cran.r-project.org/web/packages/locfit/locfit.pdf>.
- Ma, L., J. Song, E. Whinton, A. Zheng, T. Vorburger, and J. Zhou. 2004. "NIST Bullet Signature Measurement System for Rm (Reference Material) 8240 Standard Bullets." *Journal of Forensic Sciences* 49 (4): 649–59.
- National Research Council. 2009. *Strengthening Forensic Science in the United States: A Path Forward*. Washington, D.C.: The National Academies Press.
- Ripley, Brian, Bill Venables, Douglas M. Bates, Kurt Hornik, Albrecht Gebhardt, and David Firth. 2013. "MASS: Support Functions and Datasets for Venables and Ripley's Mass." Computer Program. <https://cran.r-project.org/web/packages/MASS/MASS.pdf>.